




Rows are input modalities and columns are ontologies. Bars show the mean $\pm$ SD across five seeds; higher is better. Bars are grouped/faceted, not stacked, because model scores are not additive. AUROC/AUPRC rank the models differently from F_{max}/S_{mini} ; see the paired metric-family figure. Macro-AUPRC values are provisional: archived tables were produced with a trapezoidal PR estimator that is upward-biased for tied scores (see docs/figure_data_integrity.md); regenerate before reporting. MLP under structure-only receives zeroed features and ignores graph edges, so it emits a constant score and acts as a null control, not a model. Its Macro-AUROC is 0.500 (chance); the ~ 0.70 Micro-AUROC reflects the term-frequency prior, not structure.