




Top row: Micro-F_{max}. Bottom row: Micro-AUROC. Identical layout and colour mapping in both rows — only the metric changes, and the model ranking inverts with it, so neither family alone supports a single 'best model' claim. Error bars show mean $\pm$ s.d. across $n = 5$ training seeds; points are seeds, open diamonds the mean. † marks the MLP structure-only constant-predictor control (Macro-AUROC 0.500 = chance). Generated by `src/plot_arc_ablations.py` from `ablations/**/test_metrics.json`.