




CAFA F_{max} and S_{min} carry 95% bootstrap intervals (1,000 paired protein resamples). Micro-AUPR and coverage are single point estimates in the results table with no replicates, so they are shown without error bars rather than with invented ones. Rows share a y-axis across ontologies so cross-ontology magnitudes are comparable. Coverage is the fraction of the 754 test proteins receiving any non-zero score; a method with high F_{max} but low coverage annotates few proteins confidently rather than annotating all of them well. Excluded methods: none.

Generated by `src/plot_benchmark.py` from `results/benchmark_metrics.csv`.