




Rows are input modality, columns are GO sub-ontology, and each panel shows the five architectures on a shared y-axis — modality is a facet rather than a marker shape, so no symbol decoding is required. Points are individual seeds; the horizontal bar is the mean with error bars show mean  $\pm$  s.d. across  $n = 5$  training seeds. † marks the MLP structure-only constant-predictor control. AUROC/AUPRC rank the models differently from  $F_{max}/S_{min}$ ; see the paired metric-family figure. Macro-AUPRC values are provisional: archived tables were produced with a trapezoidal PR estimator that is upward-biased for tied scores (see docs/figure\_data\_integrity.md); regenerate before reporting. MLP under structure-only receives zeroed features and ignores graph edges, so it emits a constant score and acts as a null control, not a model. Its Macro-AUROC is 0.500 (chance); the ~0.70 Micro-AUROC reflects the term-frequency prior, not structure. Generated by `src/plot_arc_ablations.py` from `ablations/**/test_metrics.json`.