




Rows are input modality, columns are ontology. Boxes show the median and interquartile range across $n=5$ training seeds and whiskers the full range; because five points is a thin basis for quartiles, every seed is plotted on top and the points are the evidence. † marks the MLP structure-only constant-predictor control. AUROC/AUPRC rank the models differently from $F_{max}S_{min}$: see the paired metric-family figure. Macro-AUPRC values are provisional: archived tables were produced with a trapezoidal PR estimator that is upward-biased for tied scores (see docs/figure_data_integrity.md); regenerate before reporting. MLP under structure-only receives zeroed features and ignores graph edges, so it emits a constant score and acts as a null control, not a model. Its Macro-AUROC is 0.500 (chance); the ~0.70 Micro-AUROC reflects the term-frequency prior, not structure. Generated by src/plot_arc_ablations.py from ablations/**test_metrics.json.