




Error bars show mean $\pm$ s.d. across $n = 5$ training seeds. F_{max}/S_{min} favour graph-aware Hybrid/Hybrid_JK. Under structure-only, MLP/GAT/GCN reach exactly the predict-nothing S_{min} (the test set's total IC), while Hybrid/Hybrid_JK saturate above the 0.99 threshold-scan bound and predict everything; both are degenerate, not tuned trade-offs. Generated by `src/plot_arc_ablations.py` from `ablations/**/test_metrics.json` + `logs/arc_ablation_*.out`.