




Error bars show mean  $\pm$  s.d. across  $n = 5$  training seeds.  $F_{max}/S_{min}$  favour graph-aware Hybrid/Hybrid\_JK. Under structure-only, MLP/GAT/GCN reach exactly the predict-nothing  $S_{min}$  (the test set's total IC), while Hybrid/Hybrid\_JK saturate above the 0.99 threshold-scan bound and predict everything; both are degenerate, not tuned trade-offs. Generated by `src/plot_arc_ablations.py` from `ablations/**/test_metrics.json` + `logs/arc_ablation_*.out`.