









Error bars show mean $\pm$ s.d. across $n = 5$ training seeds. Metric families are on separate axes with independent y-limits — AUROC occupies a much higher range and would flatten the $F_{max}/AUPR$ family on a shared scale. 'micro f1 at validation threshold' is the honest operating point: the decision threshold is selected on validation and applied unchanged to test. 'micro fmax' scans thresholds on the test set itself and is therefore an optimistic upper bound; the gap between the two is the cost of threshold transfer. The former 'micro fmax diagnostic' column is dropped here — it is numerically identical to 'micro fmax' (same test-side threshold scan), not an independent check.