




Error bars show mean $\pm$ s.d. across $n = 5$ training seeds. All three panels share one y-axis (0 to 99.6, the global maximum plus 10% headroom) so cross-ontology S_{min} magnitudes are directly comparable. Generated by `src/plot_arc_tuning.py` from `arc_tuning_cafa/test_evaluation/*/per_seed_metrics.csv`.