




Seeded random search, 40 trials (search seed 20260714); learning rate and weight decay log-uniform over $[1e-5, 3e-3]$ and $[1e-7, 1e-2]$, remaining hyperparameters categorical. Not a grid and not Bayesian. Marginal effect of each tuned hyperparameter on validation Macro- F_{max} , one point per trial. Shared y-limits across all panels. Generated by `src/plot_arc_tuning.py` from `hybrid_search/trial_*/%/config.json`.